

Stock Market Analysis for Tech Stocks

In this project, we'll analyse data from the stock market for some technology stocks.

Again, we'll use Pandas to extract and analyse the information, visualise it, and look at different ways to analyse the risk of a stock, based on its performance history.

Here are the questions we'll try to answer:

- What was the change in a stock's price over time?
- What was the daily return average of a stock?
- What was the moving average of various stocks?
- What was the correlation between daily returns of different stocks?
- How much value do we put at risk by investing in a particular stock?
- How can we attempt to predict future stock behaviour?

```
In [4]: #Python Data Analysis imports
import pandas as pd
from pandas import Series, DataFrame
import numpy as np

#Visualisation imports
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid')
matplotlib inline

#To grab stock data
from pandas.io.data import DataReader
from datetime import datetime

#To handle floats in Python 2
from __future__ import division

Users/sajal/Library/Enthought/Canopy_64bit/User/lib/python2.7/site-packages/pandas/io/data.py:35: FutureWarning:
The pandas.io.data module is moved to a separate package (pandas-datareader) and will be removed from pandas in a future version.
After installing the pandas-datareader package (https://github.com/pydata/pandas-datareader), you can change the import "from pandas.io data" to "from pandas_datareader import data, w" to "from pandas_datareader import data, w"
FutureWarning)
```

We're going to analyse some tech stocks, and it seems like a good idea to look at their performance over the last year. We can create a list with the stock names, for future looping.

```
In [29]: #We're going to analyse stock info for Apple, Google, Microsoft, and Amazon
tech_list = ['AAPL', 'GOOG', 'MSFT', 'AMZN', 'YHOO']
```

```
In [30]: #Setting the end date to today
end = datetime.now()
```

```
In [30]: #Start date set to 1 year back
start = datetime(end.year-1, end.month, end.day)
```

```
In [32]: #Using Yahoo Finance to grab the stock data
for stock in tech_list:
    globals()[stock] = DataReader(stock, 'yahoo', start, end) #The globals method sets the stock name to a global variable
```

Thanks to the globals method, Apple's stock data will be stored in the AAPL global variable dataframe. Let's see if that worked.

```
In [33]: AAPL.head()
```

```
Out[33]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2015-09-23	113.629991	114.750001	113.300003	114.300000	3976700	111.808896
2015-09-24	113.250000	114.200000	112.370003	115.000000	50219500	112.582660
2015-09-25	116.440002	116.690002	114.019997	114.709999	56519000	112.588730
2015-09-28	113.849998	114.570000	112.440002	112.440002	52109000	110.086252
2015-09-29	112.830002	113.510002	107.860001	109.059999	73965400	106.777002

```
In [17]: #Basic stats for Apple's Stock
AAPL.describe()
```

```
Out[17]:
```

	Open	High	Low	Close	Volume	Adj Close
count	253.000000	253.000000	253.000000	253.000000	2.530000e+02	253.000000
mean	104.829491	105.775100	103.861146	104.858498	4.179762e+07	103.707287
std	8.079779	8.110292	8.039386	8.079594	1.169642e+07	7.736402
min	90.000000	91.650000	88.470001	90.130000	1.204642e+07	89.953242
25%	97.300000	98.200999	96.500002	97.139999	2.844520e+07	96.346095
50%	105.519997	106.300998	104.679997	105.790000	3.695700e+07	104.701896
75%	110.629997	111.769997	109.410004	110.779989	4.896700e+07	109.220001
max	123.129997	123.820000	121.820003	122.570000	1.333950e+08	120.001314

And that easily, we can make out what the stock's minimum, maximum, and average price was for the last year.

```
In [28]: #Some basic info about the dataframe
AAPL.info()
```

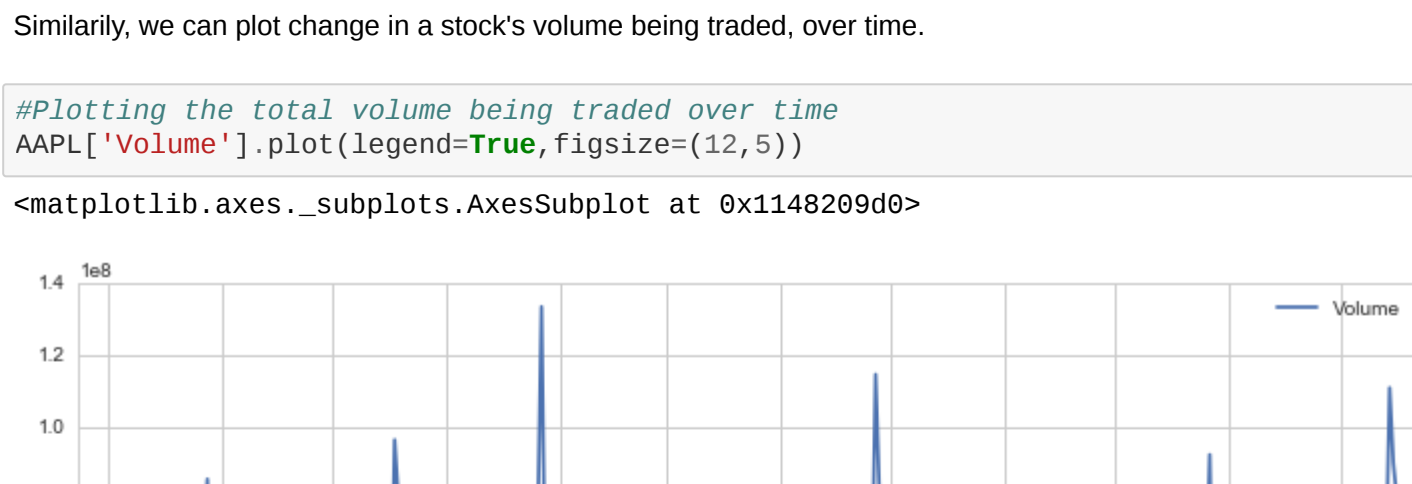
```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 253 entries, 2015-09-23 to 2016-08-22
Data columns (total 6 columns):
Open      253 non-null float64
High      253 non-null float64
Low       253 non-null float64
Close     253 non-null float64
Volume    253 non-null float64
Adj Close 253 non-null float64
dtypes: float64(5), int64(1)
memory usage: 13.8 KB
```

No missing info in the dataframe above, so we can go about our business.

What's the change in stock's price over time?

```
In [27]: #Plotting the stock's adjusted closing price using pandas
AAPL['Adj Close'].plot(legend=True, figsize=(12,5))
```

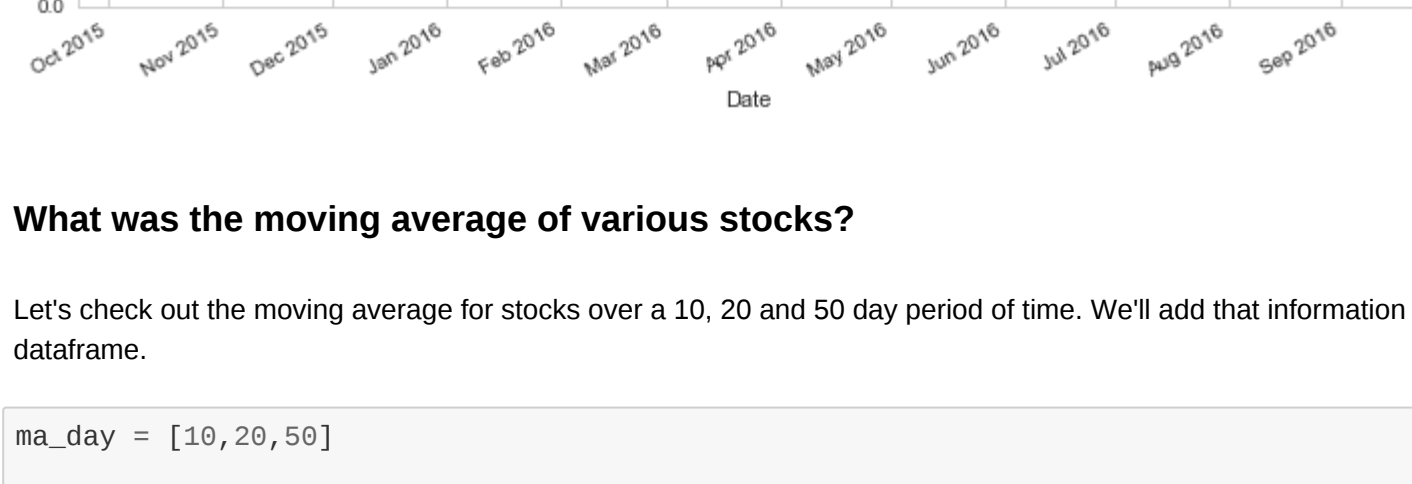
```
Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x11463ef50>
```



Similarly, we can plot change in a stock's volume being traded, over time.

```
In [34]: #Plotting the total volume being traded over time
AAPL['Volume'].plot(legend=True, figsize=(12,5))
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x1146289d8>
```



What was the moving average of various stocks?

Let's check out the moving average for stocks over a 10, 20 and 50 day period of time. We'll add that information to the stock's dataframe.

```
In [42]: ma_day = [10,20,50]
```

```
for ma in ma_day:
    column_name = "MA for %s days" %(str(ma))
    AAPL[column_name] = AAPL['Adj Close'].rolling(window=ma, center=False).mean()
```

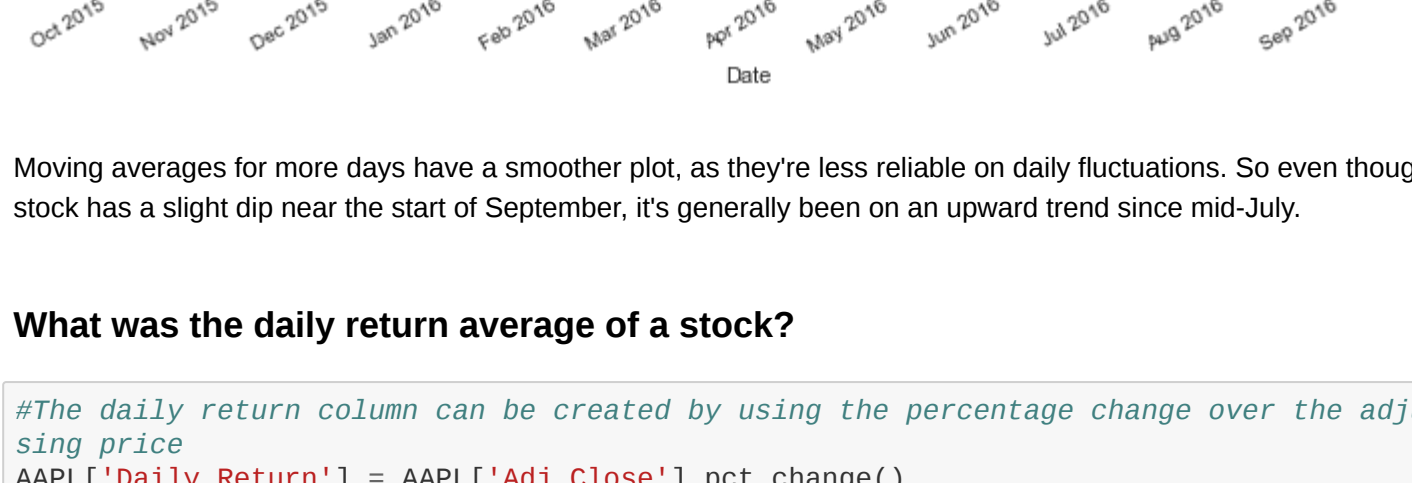
```
In [44]: AAPL.tail()
```

```
Out[44]:
```

	Open	High	Low	Close	Volume	Adj Close	MA for 10 days	MA for 20 days	MA for 50 days
Date									
2016-09-16	115.120003	116.129997	114.040001	114.919998	79886900	114.919998	108.808999	108.1500	104.962706
2016-09-19	115.190002	116.180000	113.350000	113.580002	47022300	113.580002	109.383999	108.3610	105.341124
2016-09-20	113.050003	114.120003	112.510002	113.570000	34514300	113.570000	109.980999	108.6140	105.683375
2016-09-21	113.849998	113.989998	112.440002	113.550003	36003200	113.550003	110.499999	108.8490	106.016473
2016-09-22	114.349998	114.840002	114.000000	114.620003	31017100	114.620003	111.410000	109.1785	106.301911

```
In [45]: AAPL[['Adj Close', 'MA for 10 days', 'MA for 20 days', 'MA for 50 days']].plot(subplots=False, figsize=(12,5))
```

```
Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x1146c39b0>
```



Moving averages for more days have a smoother plot, as they're less reliable on daily fluctuations. So even though, Apple's stock has a slight dip near the start of September, it's generally been on an upward trend since mid-July.

What was the daily return average of a stock?

```
In [46]: #The daily return column can be created by using the percentage change over the adjusted closing price
AAPL['Daily Return'] = AAPL['Adj Close'].pct_change()
```

```
Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x1165031d0>
```

```
In [49]: AAPL['Daily Return'].tail()
```

```
Out[49]:
```

Date	Daily Return
2016-09-16	+0.985624
2016-09-19	+0.811668
2016-09-20	-0.088088
2016-09-21	+0.088176
2016-09-22	-0.009423
Name: Daily Return, dtype: float64	

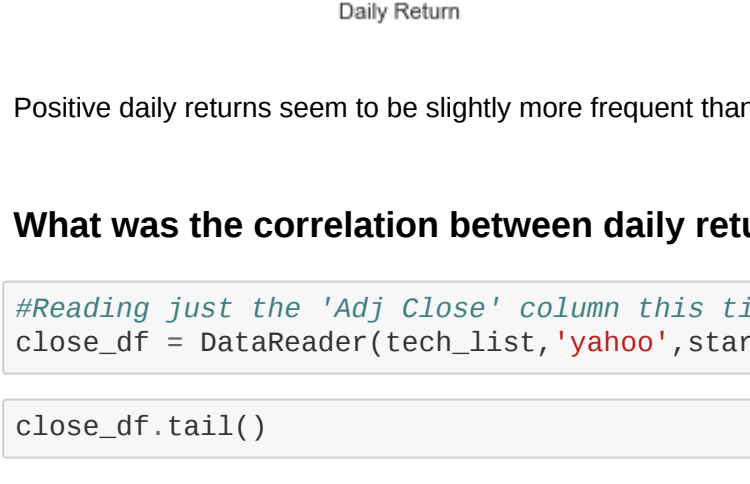
```
In [58]: #Plotting the daily return
AAPL['Daily Return'].plot(figsize=(14,5), legend=True, linestyle='--', marker='o')
```

```
Out[58]: <matplotlib.axes._subplots.AxesSubplot at 0x1156cc0b0>
```



```
In [56]: sns.distplot(AAPL['Daily Return'].dropna(), bins=100, color='red')
```

```
Out[56]: <matplotlib.axes._subplots.AxesSubplot at 0x116503aad0>
```



Positive daily returns seem to be slightly more frequent than negative returns for Apple.

What was the correlation between daily returns of different stocks?

```
In [60]: #Reading just the 'Adj Close' column of this time
close_df = DataReader(tech_list, 'yahoo', start, end)['Adj Close']
```

```
Out[60]: <matplotlib.axes._subplots.AxesSubplot at 0x1165031d0>
```

```
In [69]: close_df.tail()
```

```
Out[69]:
```

	AAPL	AMZN	GOOG	MSFT	YHOO
Date					
2016-09-16	114.919998	778.320002	768.880006	67.280000	43.669998
2016-09-19	113.580002	775.099976	765.700012	66.300000	43.189999
2016-09-20	113.570000	780.219971	771.409973	66.810001	42.790001
2016-09-21	113.550000	789.739990	776.219971	67.759999	44.139999
2016-09-22	114.620003	804.700012	787.210022	67.820000	44.150002

Everything works as expected.

Just as we did earlier, we can use Pandas' pct_change method to get the daily returns of our stocks.

```
In [66]: rets_df = close_df.pct_change()
```

```
Out[66]: <matplotlib.axes._subplots.AxesSubplot at 0x1165031d0>
```

```
In [68]: rets_df.tail()
```

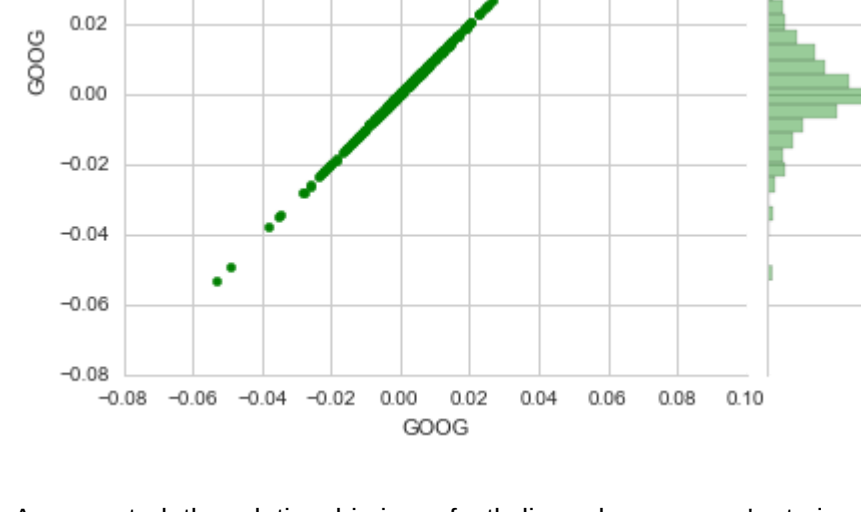
```
Out[68]:
```

	AAPL	AMZN	GOOG	MSFT	YHOO
Date					
2016-09-16	-0.005624	0.011472	-0.003732	0.001049	-0.007274
2016-09-19	-0.011660	-0.004359	-0.004136	-0.005960	-0.009992
2016-09-20	-0.000008	0.006668	0.007487	-0.002008	-0.002601
2016-09-21	0.000176	0.012002	0.002325	0.018722	0.001544
2016-09-22	0.009423	0.018943	0.014158	0.003029	0.002227

Let's try creating a scatterplot to visualise any correlations between different stocks. First we'll visualise a scatterplot for the relationship between the daily return of a stock to itself.

```
In [71]: sns.jointplot('GOOG', 'GOOG', rets_df, kind='scatter', color='green')
```

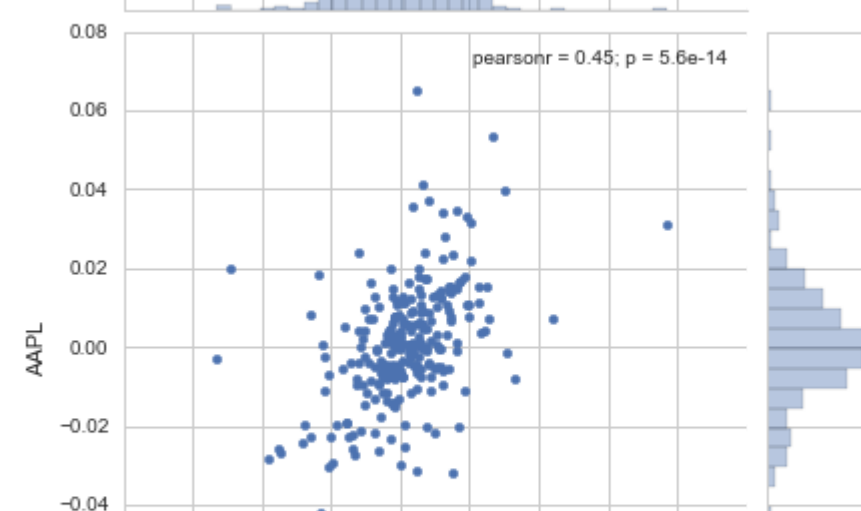
```
Out[71]: <seaborn.axisgrid.JointGrid at 0x1165031d0>
```



As expected, the relationship is perfectly linear because we're trying to correlate something with itself. Now, let's check out the relationship between 'GOOG' and 'AAPL' daily returns.

```
In [78]: sns.jointplot('GOOG', 'AAPL', rets_df, kind='scatter')
```

```
Out[78]: <seaborn.axisgrid.JointGrid at 0x11a321290>
```




There seems to be a minor correlation between the two stocks, looking at the figure above. The Pearson R Correlation Coefficient value of 0.45 echoes that sentiment.

But what about other combinations of stocks?

```
In [80]: sns.pairplot(rets_df.dropna(),)
```


```
Out[80]: <seaborn.axisgrid.PairGrid at 0x11a9ba710>
```



Quick and dirty overview visualisation of the scatterplots and histograms of daily returns of our stocks. To see the actual numbers for the correlation coefficients, we can use seaborn's corplot method.

```
In [87]: sns.corplot(rets_df.dropna(), annot=True)
```

```
Out[87]: <matplotlib.axes._subplots.AxesSubplot at 0x127a67690>
```



Google and Microsoft seem to have the highest correlation. But another interesting thing to note is that all tech companies that we explored are positively correlated.


How much value do we put at risk by investing in a particular stock?

A basic way to quantify risk is to compare the expected return (which can be the mean of the stock's daily returns) with the standard deviation of the daily returns.

```
In [90]: rets = rets_df.dropna()
```

```
In [181]: plt.figure(figsize=(8,5))
plt.scatter(rets.mean(),rets.std(),s=25)
plt.xlabel('Expected Return')
plt.ylabel('Risk')
```

```
#For adding annotations in the scatterplot
for label,x,y in zip(rets.columns,rets.mean(),rets.std()):
    plt.annotate(
        label,
        xy=(x,y),xytext=(-120,20),
        textcoords = 'offset points', ha = 'right', va = 'bottom',
        arrowprops = dict(arrowstyle='>', connectionstyle = 'arc3,rad=-0.5'))
```



We'd want a stock to have a high expected return and a low risk. Google and Microsoft seem to be the safe options for that. Meanwhile, Yahoo and Amazon stocks have higher expected returns, but also have a higher risk.

Value at Risk

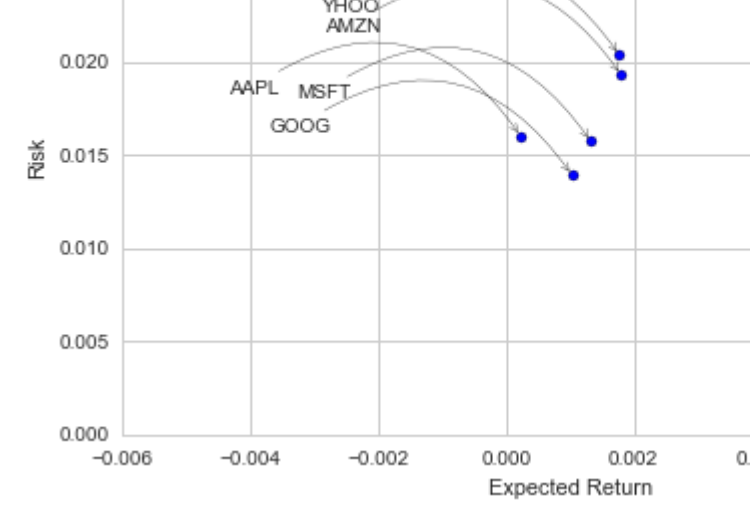
We can treat Value at risk as the amount of money we could expect to lose for a given confidence interval. We'll use the Bootstrap method and the Monte Carlo method to extract this value.

Bootstrap Method

Using this method, we calculate the empirical quantiles from a histogram of daily returns. The quantiles help us define our confidence interval.

```
In [182]: sns.distplot(AAPL['Daily Return'].dropna(), bins=100, color='purple')
```

```
Out[182]: <matplotlib.axes._subplots.AxesSubplot at 0x11534f0d0>
```



To recap, our histogram for Apple's stock looked like the above. And our daily returns dataframe looked like:

```
In [183]: rets.head()
```

```
Out[183]:
```

	AAPL	AMZN	GOOG	MSFT	YHOO
Date					
2015-09-24	0.005848	-0.004328	0.005527	0.000932	-0.013450
2015-09-25	-0.002522	-0.017799	-0.022100	0.000063	-0.007157
2015-09-28	-0.030789	-0.038512	-0.027910	-0.014793	-0.052523
2015-09-29	-0.030061	-0.015851	0.000134	0.003465	-0.029113
2015-09-30	0.011370	0.011881	0.022608	0.018877	0.023001

```
In [187]: #Using Pandas's built in quantile method
rets['AAPL'].quantile(0.95)
```

```
Out[187]: -0.825722813451247724
```

The 0.05 empirical quantile of daily returns is -0.8257. This means that with 95% confidence, the worst daily loss will not exceed 2.57% (of the investment).

How can we attempt to predict future stock behaviour?

Monte Carlo Method

Check out this [link](#) for more info on the Monte Carlo method. In short: in this method, we run simulations to predict the future many times, and aggregate the results in the end for some quantifiable value.

```
In [188]: days = 365
delta_t
dt = 1/365
mu = rets.mean()['GOOG']
sigma = rets.std()['GOOG']
```

```
In [194]: #Function takes in stock price, number of days to run, mean and standard deviation values
def stock_monte_carlo(start_price,days,mu,sigma):
    price=np.zeros(days)
    price[0] = start_price
    shock = np.zeros(days)
    drift = np.zeros(days)
    for x in xrange(1,days):
        #Shock and drift formulas taken from the Monte Carlo formula
        shock[x] = np.random.normal(loc=mu*dt,scale=sigma*np.sqrt(dt))
        drift[x] = mu * dt
        #New price = Old price + Old price*(shock+drift)
        price[x] = price[x-1] + (price[x-1] * (drift[x]+shock[x]))
    return price
```

We're going to run the simulation of Google stocks. Let's check out the opening value of the stock.

```
In [190]: GOOG.head()
```

```
Out[190]:
```

	Open	High	Low	Close	Volume	Adj Close
Date						
2015-09-23	622.949988	628.929993	620.000000	622.359985	1470500	622.359985
2015-09-24	616.640015	627.320007	612.400024	625.799988	2240100	625.799988
2015-09-25	629.770020	629.770020	611.000000	611.969971	2174000	611.969971
2015-09-28	630.340027	614.604980	589.380005	594.890015	3127700	594.890015
2015-09-29	597.280029	605.000000	590.219971	594.969971	2309500	594.969971

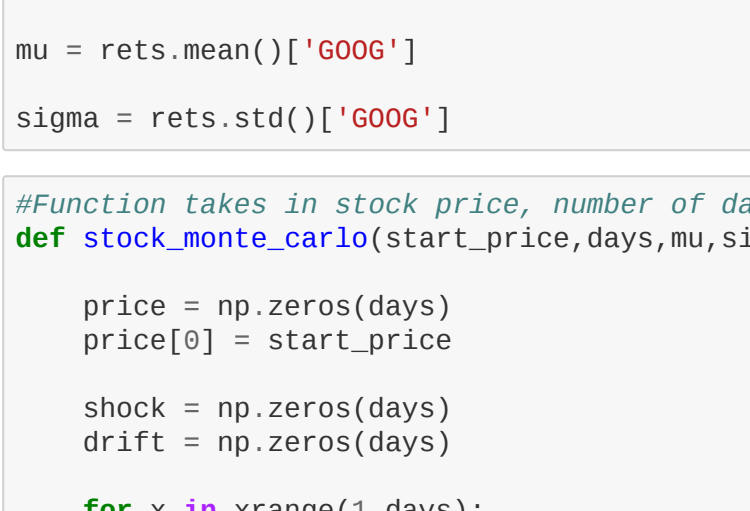
Let's do a simulation of 100 runs, and plot them.

```
In [197]: start_price = 622.849 #Taken from above
```

```
for run in xrange(100):
    plt.plot(stock_monte_carlo(start_price,days,mu,sigma))
```

```
plt.xlabel('Days')
plt.ylabel('Price')
plt.title('Monte Carlo Analysis for Google')
```

```
Out[197]: <matplotlib.text.Text at 0x11b53d0d0>
```



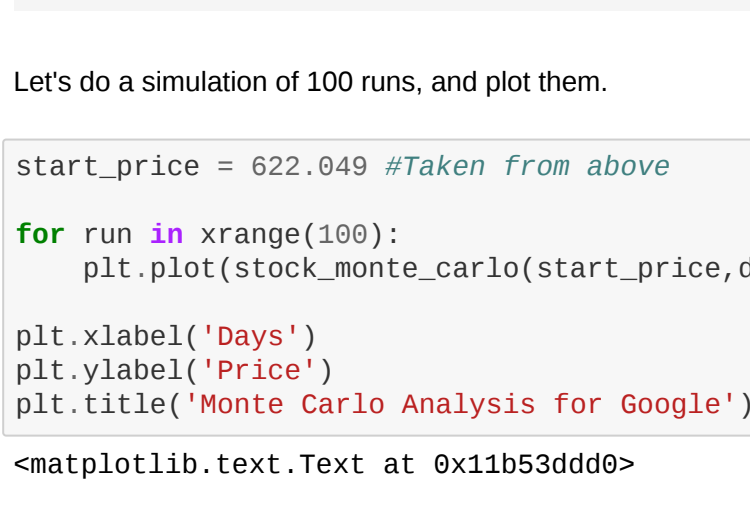
```
In [199]: runs = 10000
```

```
simulations = np.zeros(runs)
for run in xrange(runs):
    simulations[run] = stock_monte_carlo(start_price,days,mu,sigma)[days-1]
```

```
In [283]: q = np.percentile(simulations,1)
```

```
plt.hist(simulations,bins=200)
plt.figtext(0.6,0.8,"Mean Final price: %1.2f" %start_price)
plt.figtext(0.6,0.8,"Var(0.99): %1.2f" % (start_price - q))
plt.figtext(0.15,0.6,"q(0.99): %1.2f" % q)
plt.axvline(x=q, linewidth=4, color='r')
plt.title("u'Final Price distribution for Google Stock after %s days" %days, weight='bold')
```

```
Out[283]: <matplotlib.text.Text at 0x127a7e1c0>
```



We can infer from this that, Google's stock is pretty stable. The starting price that we had was USD622.05, and the average final price over 10,000 runs was USD623.36.

The red line indicates the value at risk of stock at the desired confidence interval. For every stock, we'd be risking USD18.38, 99% of the time.